

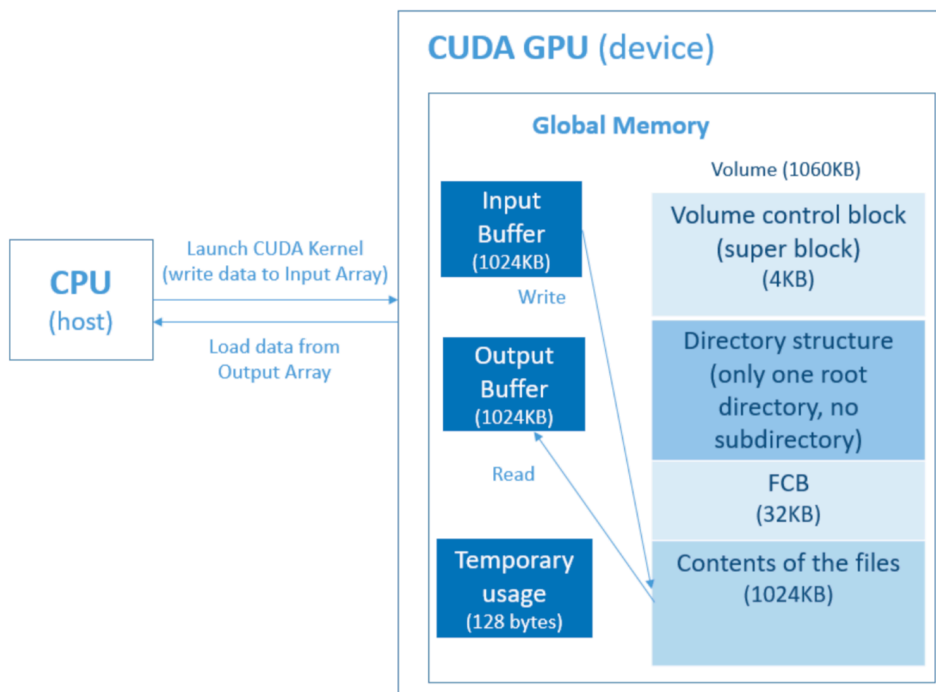
# CSC3150 Assignment4 Report

## Introduction

In this assignment, a simple mechanism of file system management was implemented via GPU's memory with single thread. There are five file operations implemented: *fs\_open*, *fs\_write*, *fs\_read*, *fs\_gsys(RM)* and *fs\_gsys(LS\_D/LS\_S)*. The source code of this file system implementation has been attached with this report, which is in the file "*file\_system.cu*".

Specification:

- The global memory was taken as a volume
- Contiguous allocation was adopted
- The size of volume - 1060 KB
- The total size of files - 1024 KB
- The maximum number of file - 1024
- The maximum size of a file - 1 KB (1024 bytes)
- The maximum size of a file name - 20 bytes (file name ends with '\0')
- FCB size - 32 bytes
- FCB entries - 32 KB / 32 bytes = 1024
- Storage block size - 32 bytes



## Program design

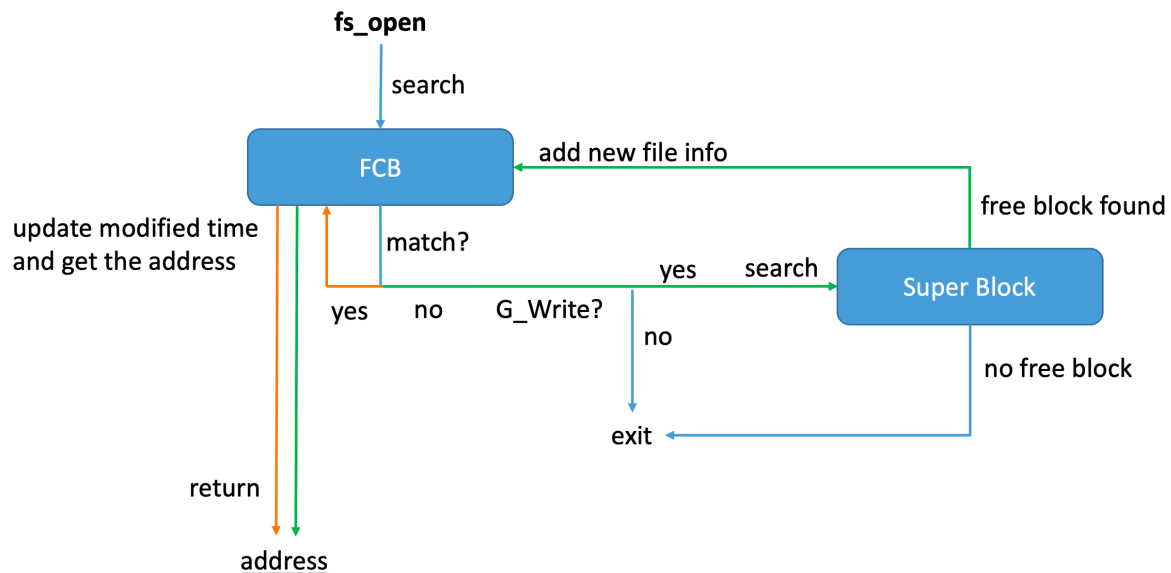
### Initialization of the super block

- The super block was initialized in the function *fs\_init*:

```
// initialize the Super Block
for (int i = 0; i < SUPERBLOCK_SIZE; ++ i){
    volume[i] = 255; // 255 means 1111 1111
}
```

where 255 means 1111 1111 in binary, since each uchar could represent a 8-bit bit map for 8 blocks.

## fs\_open



As the logic graph above shows, when the function `fs_open()` is called, the following steps was processed:

1. Search for the FCB entry based on the input file name:

```
// serch the file info in FCB
for (i = fs->SUPERBLOCK_SIZE; i < fs->FILE_BASE_ADDRESS; i += fs->FCB_SIZE){
    is_match = true;
    // comparing the file name
    for (j = 0; j < fs->MAX_FILENAME_SIZE; ++j) {
        if (s[j] != fs->volume[i + j]) {
            is_match = false;
            break;
        }
    }
    if (is_match) {
        index = i; // The FCB entry address of the target file
        break;
    }
}
```

2. Return the physical address if file name matched:

```

if (index != -1) {    // the file exists
    // update the modified time
    fs->volume[index + 31] = gtime % 256;
    fs->volume[index + 30] = (gtime >> 8) % 256;
    // get address from FCB
    u32 address = fs->volume[index+23] + (fs->volume[index+22]<<8) + (fs->volume[index+21]<<16) + (fs->volume[index+20]<<24);
    return address;
}

```

3. If no matched file name and G\_Write is input op, then search in Super Block:

```

for (i = 0; i < fs->SUPERBLOCK_SIZE; i += 4){
    if (fs->volume[i] == 255){ // every uchar represents 8 bits(block)
        fs->volume[i] = 127;
        index = i * 8;        // #start block of the file
        break;
    }
} // end for

```

4. If there is a free block, calculate the entry address of FCB and the file information:

```

base = fs->SUPERBLOCK_SIZE + index;
start = fs->FILE_BASE_ADDRESS + index * (fs->STORAGE_BLOCK_SIZE); // the
physical address of the file in byte
size = 0; // the size of the file in in byte
created = gtime; // get current time as created time
modified = gtime; // get current time as modified time

```

5. Store the information to FCB and return the physical address of the file:

```

// store the file name
for (j = 0; j < fs->MAX_FILENAME_SIZE; ++j) {
    fs->volume[base + j] = s[j];
    if (s[j] == '\0') break;
}

// store the physical start address of the file (4B)
fs->volume[base + 23] = start % 256;
fs->volume[base + 22] = (start >> 8) % 256;
fs->volume[base + 21] = (start >> 16) % 256;
fs->volume[base + 20] = (start >> 24) % 256;

// store the size of the file (4B)
fs->volume[base + 27] = size % 256;
fs->volume[base + 26] = (size >> 8) % 256;

```

```

// store the created time
fs->volume[base + 29] = created % 256;
fs->volume[base + 28] = (created >> 8) % 256;

// store the modified time
fs->volume[base + 31] = modified % 256;
fs->volume[base + 30] = (modified >> 8) % 256;

return start;

```

## fs\_read

For `fs_read()`, simply read the data from physical storage based on the input file pointer and copy the data to the output buffer:

```

__device__ void fs_read(FileSystem *fs, uchar *output, u32 size, u32 fp)
{
    for (int i = 0; i < size; ++ i) {
        output[i] = (char)fs->volume[fp++];
    }
}

```

## fs\_write

For file writing:

1. Increment the timing since the file is going to be modified:

```
gtime++;
```

2. Based on the input file pointer, calculate the corresponding entry address of FCB:

```

u32 index, base, file_size;
index = (fp - fs->FILE_BASE_ADDRESS) >> 10; // the index of the entry of
FCB
base = fs->SUPERBLOCK_SIZE + index * fs->FCB_SIZE; // address of the
entry

```

3. Get the size of the old file for later storage cleanup use:

```

file_size = fs->volume[base+27] + (fs->volume[base+26]<<8) + (fs-
>volume[base+25]<<16) + (fs->volume[base+24]<<24);

```

4. Update the size of the file and the modified time:

```
// update the modified time
fs->volume[base+31] = gtime % 256;
fs->volume[base+30] = (gtime >> 8) % 256;

// update the size of file
fs->volume[base + 27] = size % 256;
fs->volume[base + 26] = (size >> 8) % 256;
```

5. Write the physical storage, update the super block, and cleanup the contents of the old file based on the input file pointer:

```
// write the physical storage
for (int i = 0; i < size; ++ i){
    // update the super block everytime writing a new block
    if (i % 32 == 0) {
        int block_order = i / 32;
        uchar target = fs->volume[index * 4 + block_order / 8];
        if (target / (1 << (7 - (block_order % 8))) % 2 == 1)
            target -= (1 << (7 - (block_order % 8))); // target block: 1
        -> 0
        fs->volume[index * 4 + block_order / 8] = target;
    }
    // write
    fs->volume[fp++] = input[i];
}

// cleanup the older contents left
if (file_size > size) {
    for (int i = fp; i < fp + file_size - size; ++i){
        fs->volume[i] = '\0';
    }
}
```

6. Return the file pointer.

### fs\_gsys(LS\_D/LS\_S)

Basically, regardless of the printing information, the logic of the implementation of LS\_D and LS\_S are the same, except that LS\_D sorts the modified time and LS\_S sorts the file size. Thus here just elaborated the implementation of LS\_D only. The insertion sort was adopt as the sorting method, which means it put the newly added elements to the right position everytime add a new element into the array.

1. Prepare the array for sorting:

```
u32 sort_arr[fs->MAX_FILE_NUM];
int arr_size = 0;
```

2. Get modified time (for LS\_S, size) information from FCB and do the sorting:

```
for (int i = 0; i < fs->FCB_ENTRIES; ++ i){
    base = fs->SUPERBLOCK_SIZE + i * fs->FCB_SIZE;
    // FCB entry is not empty
    if (fs->volume[base] != '\0'){
        entry = i;
        modified = fs->volume[base+31] + (fs->volume[base+30] << 8);
        int ptr = arr_size - 1;
        sort_arr[arr_size++] = (entry << 16) + modified; // add to the
sort array
        // insertion sort
        while (ptr >= 0 && modified < sort_arr[ptr] % (1 << 16)) {
            sort_arr[ptr + 1] = sort_arr[ptr];
            sort_arr[ptr] = (entry << 16) + modified;
            ptr--;
        } // end while
    } // end if
} // end for
```

Notice that, during the sorting, the information of the FCB entry and the modified time were translated as a *u32* data, since *u32* data type takes 4 bytes and the FCB entry and modified time take 2 bytes each. Here, the most significant two bytes represents FCB entry and the least significant two bytes represents modified time. The sorting only focus on the modified time (least significant two bytes) and the FCB entry would be extracted for locating the file name after sorting.

3. Print the sorting results:

```
printf("===sort by modified time===\n");
while (--arr_size >= 0) {
    entry = sort_arr[arr_size] >> 16;
    printf("%s\n", &fs->volume[fs->SUPERBLOCK_SIZE + entry * fs->FCB_SIZE]);
}
```

## fs\_gsys(RM)

For file removal:

1. Search for the FCB entry based on the input file name, which is the same as what in the *fs\_open()*

```
int index = -1;
for (i = fs->SUPERBLOCK_SIZE; i < fs->FILE_BASE_ADDRESS; i += fs->FCB_SIZE) {
    is_match = true;
    // comparing the file name
    for (j = 0; j < fs->MAX_FILENAME_SIZE; ++j) {
```

```

        if (s[j] != fs->volume[i + j]) {
            is_match = false;
            break;
        }
    }
    if (is_match) {
        index = i; // address of the FCB entry
        break;
    }
} // end for

```

2. If the file is found, get the address and the size of the file:

```

u32 address = fs->volume[index + 23] + (fs->volume[index + 22] << 8) +
(fs->volume[index + 21] << 16) + (fs->volume[index + 20] << 24);
u32 size = fs->volume[index + 27] + (fs->volume[index + 26] << 8);

```

3. Release the space of the file based on the address:

```

for (i = 0; i < size; ++i) {
    fs->volume[address + i] = '\0';
}

```

4. Release the corresponding FCB:

```

for (i = 0; i < fs->FCB_SIZE; ++i) {
    fs->volume[index + i] = '\0';
}

```

5. Update the super block:

```

index = (index - fs->SUPERBLOCK_SIZE) / fs->FCB_SIZE;
for (i = 0; i < 4; ++i) {
    fs->volume[index * 4 + i] = 255;
}

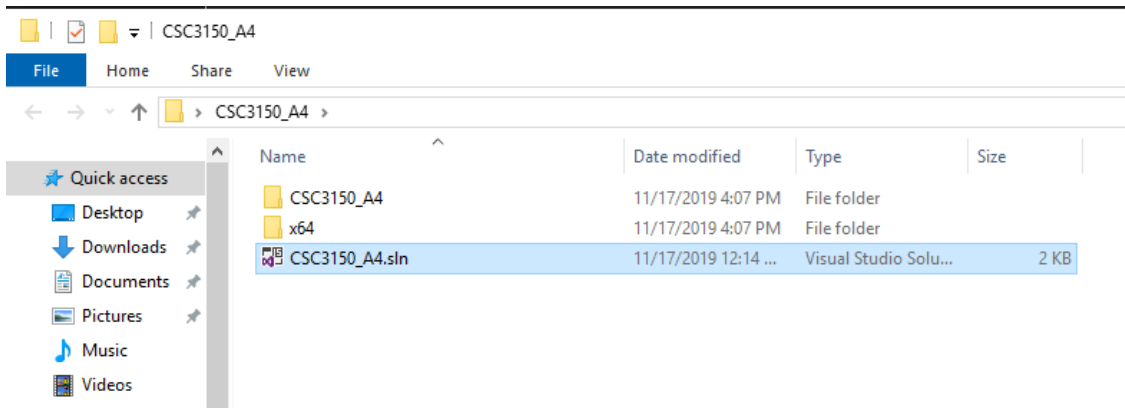
```

## Running environment

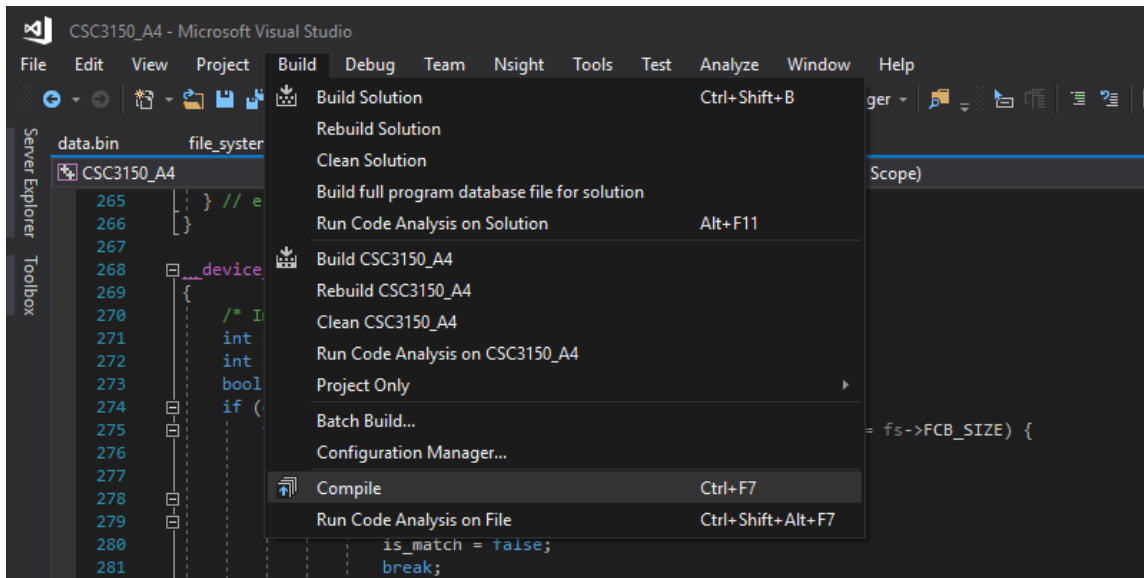
- Windows 10 Enterprise
- Microsoft Visual Studio 2017
- CUDA 9.2
- NVIDIA Geforce GTX 1060 6GB
- Compute capacity: 6.1

## Steps to run the program

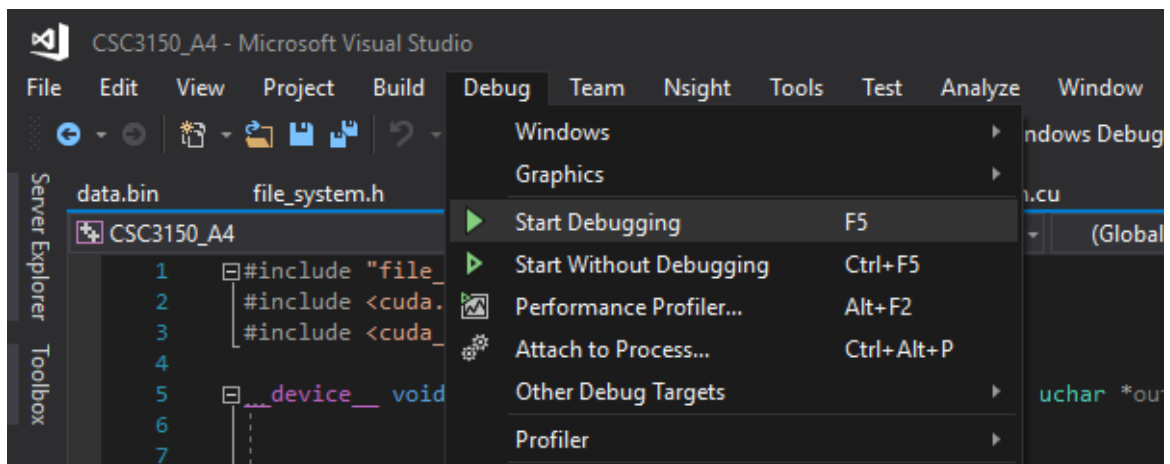
1. In the folder of "CSC3150\_A4", open "CSC3150\_A4.sln" with VS 2017:



2. In the VS 2017, compile all the file end with ".cu" by pressing 'Ctrl' + 'F7' :



3. Select "Debug" item in the menu and click "Start Debugging" to run the program (or directly press F5):



## Sample output

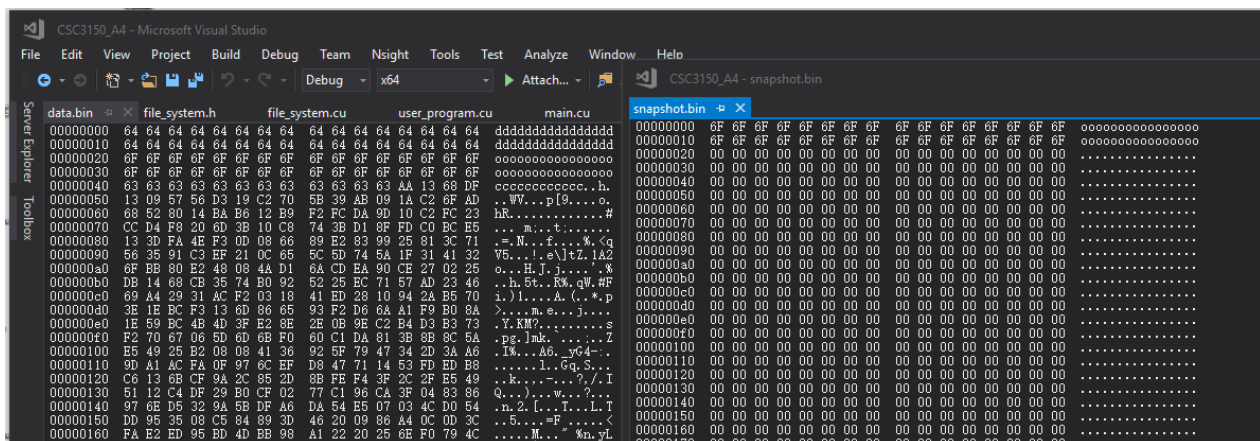
Test case 1



```
Microsoft Visual Studio Debug Console

===sort by modified time===
t.txt
b.txt
===sort by file size===
t.txt 32
b.txt 32
===sort by file size===
t.txt 32
b.txt 12
===sort by modified time===
b.txt
t.txt
===sort by file size===
b.txt 12

C:\Users\117020343\Desktop\CSC3150_A4\x64\Debug\CSC3150_A4.exe (process 17172) exited with code 0.
To automatically close the console when debugging stops, enable Tools->Options->Debugging->Automatically close the console when debugging stops.
Press any key to close this window . . .
```

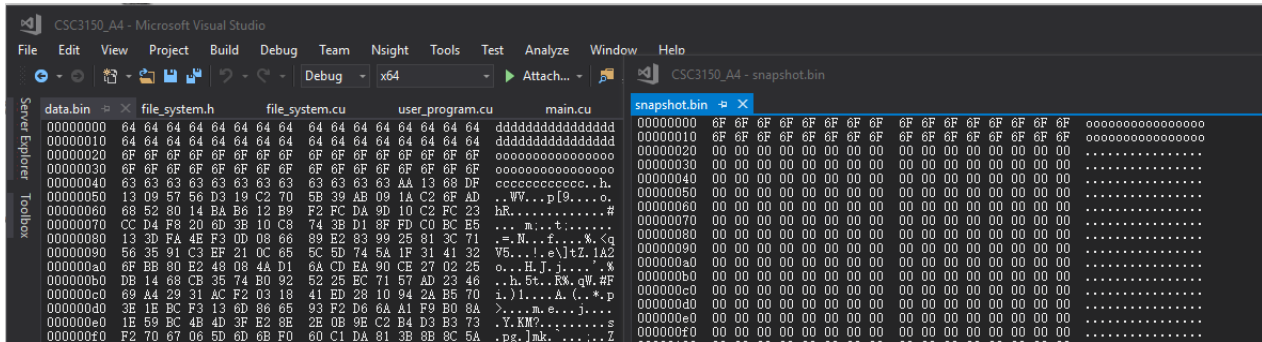


## Test case 2

```
Microsoft Visual Studio Debug Console

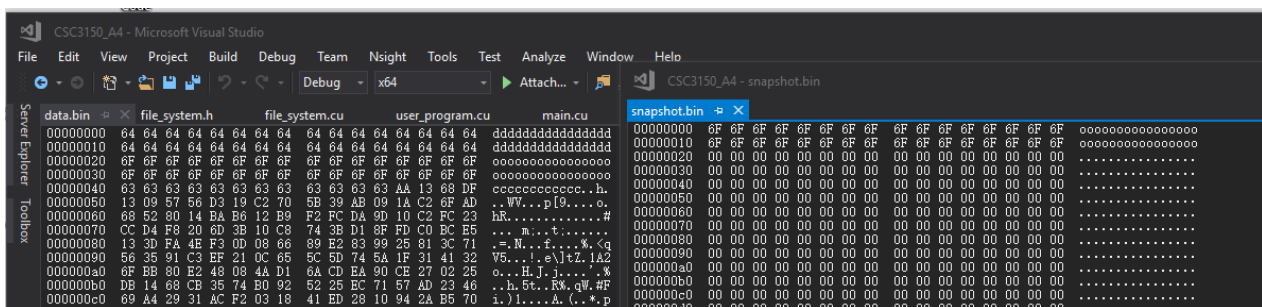
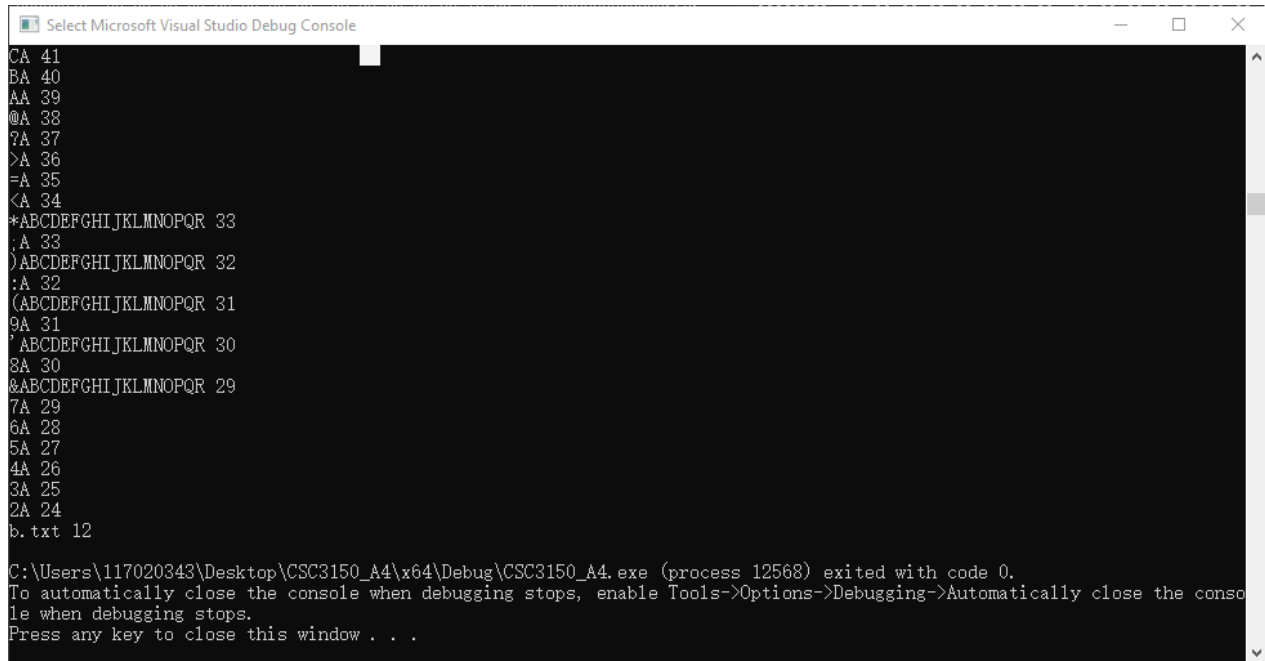
===sort by modified time===
t.txt
b.txt
===sort by file size===
t.txt 32
b.txt 32
===sort by file size===
t.txt 32
b.txt 12
===sort by modified time===
b.txt
t.txt
===sort by file size===
b.txt 12
===sort by file size===
*ABCDEF GHIJ KLMNOPQR 33
)ABCDEF GHIJ KLMNOPQR 32
(ABCDEF GHIJ KLMNOPQR 31
'ABCDEF GHIJ KLMNOPQR 30
&ABCDEF GHIJ KLMNOPQR 29
%ABCDEF GHIJ KLMNOPQR 28
$ABCDEF GHIJ KLMNOPQR 27
#ABCDEF GHIJ KLMNOPQR 26
"ABCDEF GHIJ KLMNOPQR 25
!ABCDEF GHIJ KLMNOPQR 24
b.txt 12
===sort by modified time===
*ABCDEF GHIJ KLMNOPQR
)ABCDEF GHIJ KLMNOPQR
(ABCDEF GHIJ KLMNOPQR
'ABCDEF GHIJ KLMNOPQR
&ABCDEF GHIJ KLMNOPQR
b.txt

C:\Users\117020343\Desktop\CSC3150_A4\x64\Debug\CSC3150_A4.exe (process 3152) exited with code 0.
```



## Test case 3 (part)

The output for Test case 3 has thousands of lines, here only show the last several lines.



## Encountered problems

### Problem

The given FCB size is limited to 32 bytes and each byte is stored in form of *uchar*, but each kind of information to be stored is larger than 1 bytes and inform of u32 or other types.

### Solution

To make the most of each bytes in FCB, the distribution was made carefully. For each entry of the FCB, it has 32 bytes and was specified as follows:

- File name - 20 bytes

- The physical address of the file - 4 bytes
- The size of the file - 2 bytes
- The created time - 2 bytes
- The modified time - 2 bytes.

Other information of files is not covered since it is not used in this task.

To store the information into FCB or extract the information from FCB, some transformations are required:

- For file names, since it is in form of *char* with 1 byte each, simply copy the name character by character, byte by byte.
- For the other information, they are in form of *u32*, which is 4 bytes each, so bit operations are needed. The integer was regarded as 32-bit binary number and was partitioned as several 8-bit numbers (1 byte each) to store. And these 8-bit number was combined together again when the information was extracted from the FCB.

Set the physical address of files as an example:

Store:

```
// store the physical start address of the file (4B)
fs->volume[base + 23] = addr % 256;
fs->volume[base + 22] = (addr >> 8) % 256;
fs->volume[base + 21] = (addr >> 16) % 256;
fs->volume[base + 20] = (addr >> 24) % 256;
```

Extract:

```
u32 address = fs->volume[index+23] + (fs->volume[index+22]<<8) + (fs->
volume[index+21]<<16) + (fs->volume[index+20]<<24);
```

## Problem

When the program run the test case 1, it exited unexpectedly.

## Solution

This may be caused by segment fault, which refers to the problem of memory overflow. Function *printf()* was used to keep track of the variables during program running and the output showed that some variables went wrong before the program exited. To locate the problem, more *printf()* was used to track. Finally the problem was found at the bit operation. As for the priority of the operators, "+", "-", "\*", "/", "%" are higher than "<<" and ">>", but the operations of "<<" and ">>" in my code were not put in brackets, which resulted in the problem.

## What I learned

Better understanding of the mechanism of file system

- How does each part work together: volume structure, super block, file control block (FCB), free-space management etc.

- How to allocate the file into physical storage: the implementation of contiguous allocation
- To make best of the space, some operation might needed to translate the data.
- Better understanding of the implementation of file related system call and how they work together to attain some user operations.
- ...

## Bonus

### Tree structure

To implement the directory, tree structure was adopted and a new block in the volume was taken to store the information of the tree. Since there are at most 1024 files, so totally 1024 entries for the directory and each entry takes 27 bytes with following distribution:

Entry	is_dir	id	parent_id	num_files	depth	file_name
Bytes taken	1	2	2	1	1	20

where:

- **is\_dir** is to identify whether the file is a directory or not.
- **id** is the id of the file, which is equal to the index of the entry, so id is smaller than 1024.
- **parent\_id** records the id of the parent of the file. (for root directory, its parent id was set as 1024)
- **num\_files** records how many files (including directories) contained in this directory, for non-directory, this variable is 0. This is to avoid the files under a directory exceeds 50.
- **depth** records the depth of the file in the tree structure, this is to avoid the depth of tree exceeds 3.

To keep the track of the current direcorey, a global variable `current_dir` was set and initilized as 0:

```
__device__ __managed__ u32 current_dir = 0;
```

To implement the directory, we just need to maintain the structure above and the global variable above.

### Root directory initialization

In the function `fs_init()`:

```
// dir flag
// 1 means directory file, 0 means non-directory file
fs->volume[fs->DIR_BASE_ADDRESS] = 1;

// dir id
fs->volume[fs->DIR_BASE_ADDRESS+2] = 0;
fs->volume[fs->DIR_BASE_ADDRESS+1] = 0;
```

```

// parent id
fs->volume[fs->DIR_BASE_ADDRESS+4] = 1024 % 256; // since no id would be 1024,
fs->volume[fs->DIR_BASE_ADDRESS+3] = 1024 >> 8; // set parent as 1024
represents root dir

// number of files contained
fs->volume[fs->DIR_BASE_ADDRESS+5] = 0;

// depth
fs->volume[fs->DIR_BASE_ADDRESS+6] = 0;

// name
fs->volume[fs->DIR_BASE_ADDRESS+7] = '/';
fs->volume[fs->DIR_BASE_ADDRESS+8] = '\0';

```

## fs\_gsys(MKDIR)

1. Get the depth of the current directory in the tree structure, and check whether it is smaller than 3. If is smaller than 3, continue; otherwise, exit.

```

int depth = fs->volume[fs->DIR_BASE_ADDRESS + current_dir*fs->DIR_SIZE +
6];

```

2. Search in the directory strucutre, find an empty entry:

```

for (int i = 0; i < fs->DIR_ENTRIES; ++ i){
    // find empty dir entry
    if (fs->volume[fs->DIR_BASE_ADDRESS + i*fs->DIR_SIZE] == '\0'){
        new_id = i;
        break;
    }
}

```

3. Add the information of the new directory into directory structure:

```

int base = fs->DIR_BASE_ADDRESS + new_id * fs->DIR_SIZE; // base address

// dir flag
fs->volume[base] = 1;

// dir id
fs->volume[base+2] = new_id % 256;
fs->volume[base+1] = new_id >> 8;

// parent id
fs->volume[base+4] = current_dir % 256; // not id would be 1024,

```

```

fs->volume[base+3] = current_dir >> 8; // set parent as 1024 represents
root dir

// number of files contained
fs->volume[base+5] = 0;

// depth
fs->volume[base+6] = depth + 1;

// name
for (int j = 0; j < fs->MAX_FILENAME_SIZE; ++j) {
    fs->volume[base+7+j] = s[j];
    if (s[j] == '\0') break;
}

```

4. Update the FCB and the physical storage of files:

```

// update FCB and file space
fs_open(fs, s, 1);

```

5. Update the parent directory in directory structure:

```

// increment the num_files of the parent directory
fs->volume[fs->DIR_BASE_ADDRESS+dir_index*fs->DIR_SIZE+5] += 1;

```

## fs\_gsys(CD)

Search the directory structure for the target directory, by focusing the subdirectories of the current directory. After find the id of the target directory, update the *current\_dir* so as to enter the directory:

```

int parent, base;
bool is_match;
// serch in directory
for (int i = 0; i < fs->DIR_ENTRIES; ++i){
    base = fs->DIR_BASE_ADDRESS + i * fs->DIR_SIZE;
    parent = fs->volume[base + 4] + (fs->volume[base+3] << 8);
    // only focus on subdirectory
    if (parent == current_dir){
        is_match = true;
        for (int j = 0; j < fs->MAX_FILENAME_SIZE; ++j) {
            if (s[j] != fs->volume[i + j]) {
                is_match = false;
                break;
            }
        }
    }
    // update current directory
    if (is_match) {

```

```

        current_dir = i;
        break;
    }
}
} // end for

```

## fs\_gsys(CD\_P)

```

// get the base address of current directory
int base = fs->DIR_BASE_ADDRESS + current_dir * fs->DIR_SIZE;
// get the id of parent directory from directory structure
int parent = fs->volume[base+4] + (fs->volume[base+3] << 8);
// update the current directory
current_dir = parent;

```

## fs\_gsys(PWD)

```

int print_list[3];
int size = 0;
int id = current_dir;
int base, name_address;

// get all the id through the path
while (id != 0){
    print_list[size++] = id;
    base = fs->DIR_BASE_ADDRESS + id * fs->DIR_SIZE;
    id = fs->volume[base+4] + (fs->volume[base+3] << 8); // parent id
}

// print the result
if (size == 0) printf('/'); // root directory
while (size > 0){
    id = print_list[--size];
    base = fs->DIR_BASE_ADDRESS + id * fs->DIR_SIZE;
    name_address = fs->volume[base+7];
    printf("/%s", &name_address);
}

```